

1. Which of these do we want our (binomial classification) models to produce more of?

1 / 1 point

☐ False positives

☒ True negatives

✓ **Correct**

We said "no", and the answer was "no".

☒ True positives

✓ **Correct**

We said "yes", and the answer was "yes".

☐ False negatives

2. What is the response variable?

1 / 1 point

- ☐ It is another word for the model's F1 score.
- ☐ It is the model's score, an indicator of how accurate it is.
- ☒ It is the thing you are trying to learn.
- ☐ It is the column(s) you learn from, also called a predictor variable.

✓ **Correct**

Also referred to as "y" in this course.

3. When you run H2O on the same machine as your client, where was the training data that h2o deep learning used to learn from?

1 / 1 point

- ☒ In the memory of the h2o Java process.
- ☐ In the memory of your R/Python client session.
- ☐ On a remote server.
- ☐ On local disk, in a temporary file.

✓ **Correct**

Yes, everything happens in the java processes memory space.

4. What were the three arguments we need to give, as a minimum, to the h2o deep learning function?

1 / 1 point

- ☐ Our training data, our test data, and the 80-20 split we used.
- ☒ The list of predictor variables, the response variable, and the training data.
- ☐ What to predict, which columns to use to predict from, and the test data to score on.

✓ **Correct**

Yes, "x", "y" and "train".

5. AUC stands for what?

1 / 1 point

- ☐ Awkward-udder-cow
- ☒ Area-under-curve (talking about the ROC curve)
- ☐ Always-unavailable-classifier (talking about binomial classifiers)

✓ Correct

Yes, where 1.0 means a perfect model, and 0.5 means random guessing.

6. If you try to build a deep learning model with "y" set to be a column that is an enum with exactly 2 values (cardinality of two), what will you be building?

1 / 1 point

("enum" is another way of saying a factor or categorical column)

- ☒ A binomial classifier
- ☐ A regression model
- ☐ A multinomial classifier
- ☐ None of the above: deep learning will refuse to run.

✓ Correct

Yes, binomial meaning two choices.

7. I did an experiment with 25 random train/test splits of the iris data, 20 models for each, so 500 models altogether. Based on the video, and your own experiments, which do you think is the outcome. (The % refers to the percentage that it got correct.)

1 / 1 point

- ☒ Results ranged from 85% to 100%. 85 of the 500 models scored 100%.
- ☐ Results were normally distributed around 70%.
- ☐ 400 models scored 93%, the other 100 models scored 97%.

✓ Correct

Yes. The train/test split can make a big difference.
